

Ada Diaconescu

ada.diaconescu@telecom-paris.fr



What is MQTT?

Client Server Publish / Subscribe messaging transport protocol

 Characteristics: lightweight (compared to e.g. HTTP), open, simple, easy to implement (on client side)

- Ideal for constrained environments (small code footprint & limited bandwidth)
 - E.g. Machine to Machine (M2M)
 - E.g. Internet of Things (IoT)

Homepage: https://mqtt.org



MQTT History

- Invented in 1999 by Andy Stanford-Clark (IBM) & Arlen Nipper (Arcom → Cirrus Link)
- Needed a protocol for minimal battery loss and minimal bandwidth to connect with oil pipelines via satellite
- Requirements: simple implementation, Quality of Service (QoS) data delivery,
 lightweight & bandwidth efficient, data agnostic, continuous session awareness
- Focus has shifted from proprietary embedded systems to open IoT
- MQTT acronym
 - Present: simply the protocol name
 - Past: MQ Telemetry Transport (based on IBM's MQ Series)



MQTT History (2)

- 1999: proprietary MQTT protocol invention (over TCP/IP)
- 2010: IBM releases royalty-free MQTT v3.1.0
- 2014: MQTT v3.1.1 standard approved by OASIS (open organisation for advancing standards, e.g. AMQP, SAML, ...) & ISO/IEC 20922 → in common use
- 2019: MQTT v₅ standard approved by OASIS → currently limited use
 - New features for Cloud platforms
 - More reliability for mission-critical messaging
- 2013: MQTT-SN for Sensor Networks over UDP, Zigbee, other transports...
 - Suitable for Internet of Things (IoT)

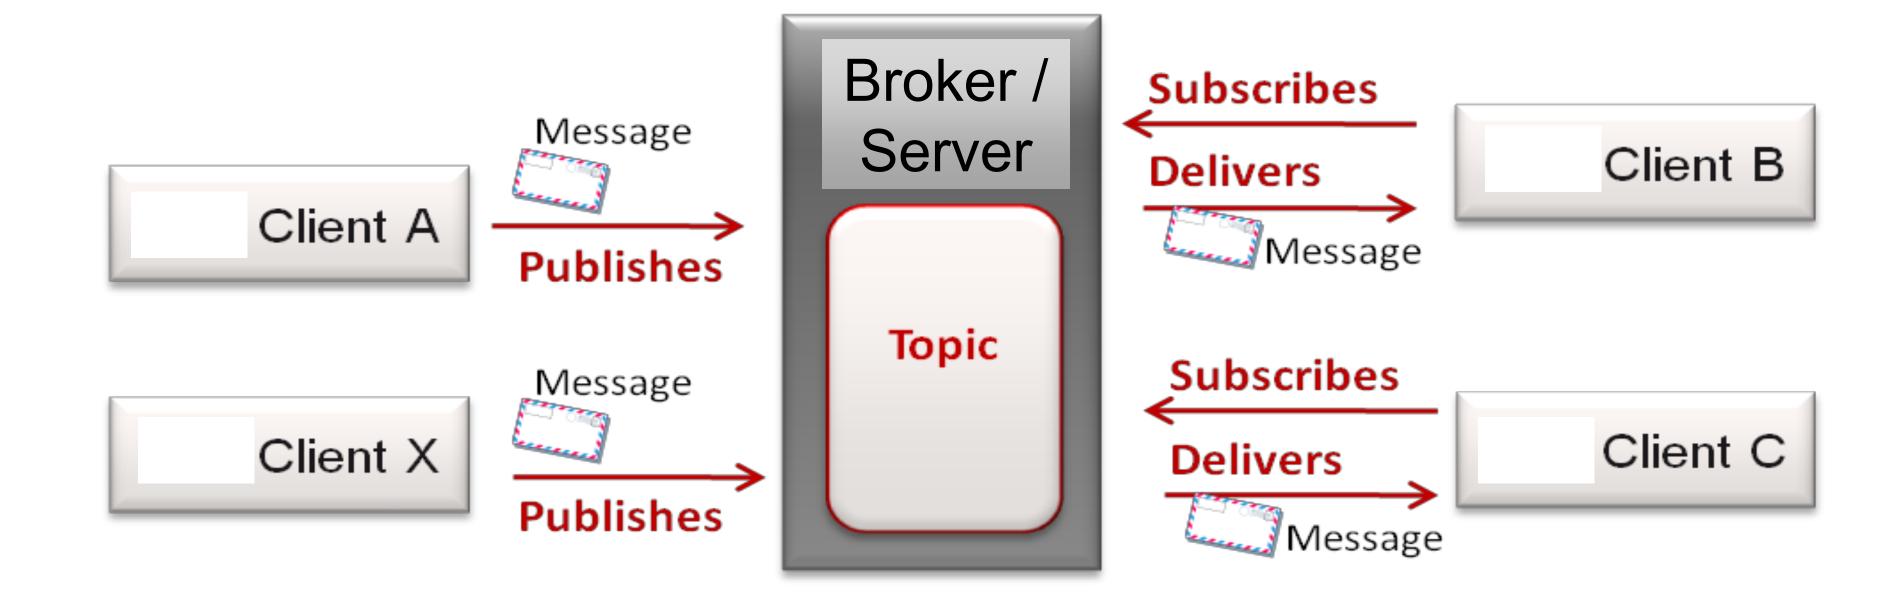


Publish/Subscribe (Pub/Sub) Communication Model

- Alternative to Client-Server model: direct communication
- Pub/Sub decouples message senders and receivers
- Model:
 - Publisher: the Client that sends the message
 - Subscriber: the Client that receives the message
 - → Publishers & Subscribers: are never in direct contact; unaware of each-other
 - Broker / Server: third party component that handles Pub/Sub client connections
 - Filter incoming messages from publishers (senders)
 - Distribute messages correctly to subscribers (receivers)



(Pub/Sub) Communication Model (2)



- Space decoupling: clients (Pub/Sub) do not need to know each-other (IP addr. & port)
- Time decoupling: clients do not need to run at the same time
- Synchronisation decoupling: clients not disrupted when sending or receiving messages



Pub/Sub example -> Car Speed Meter

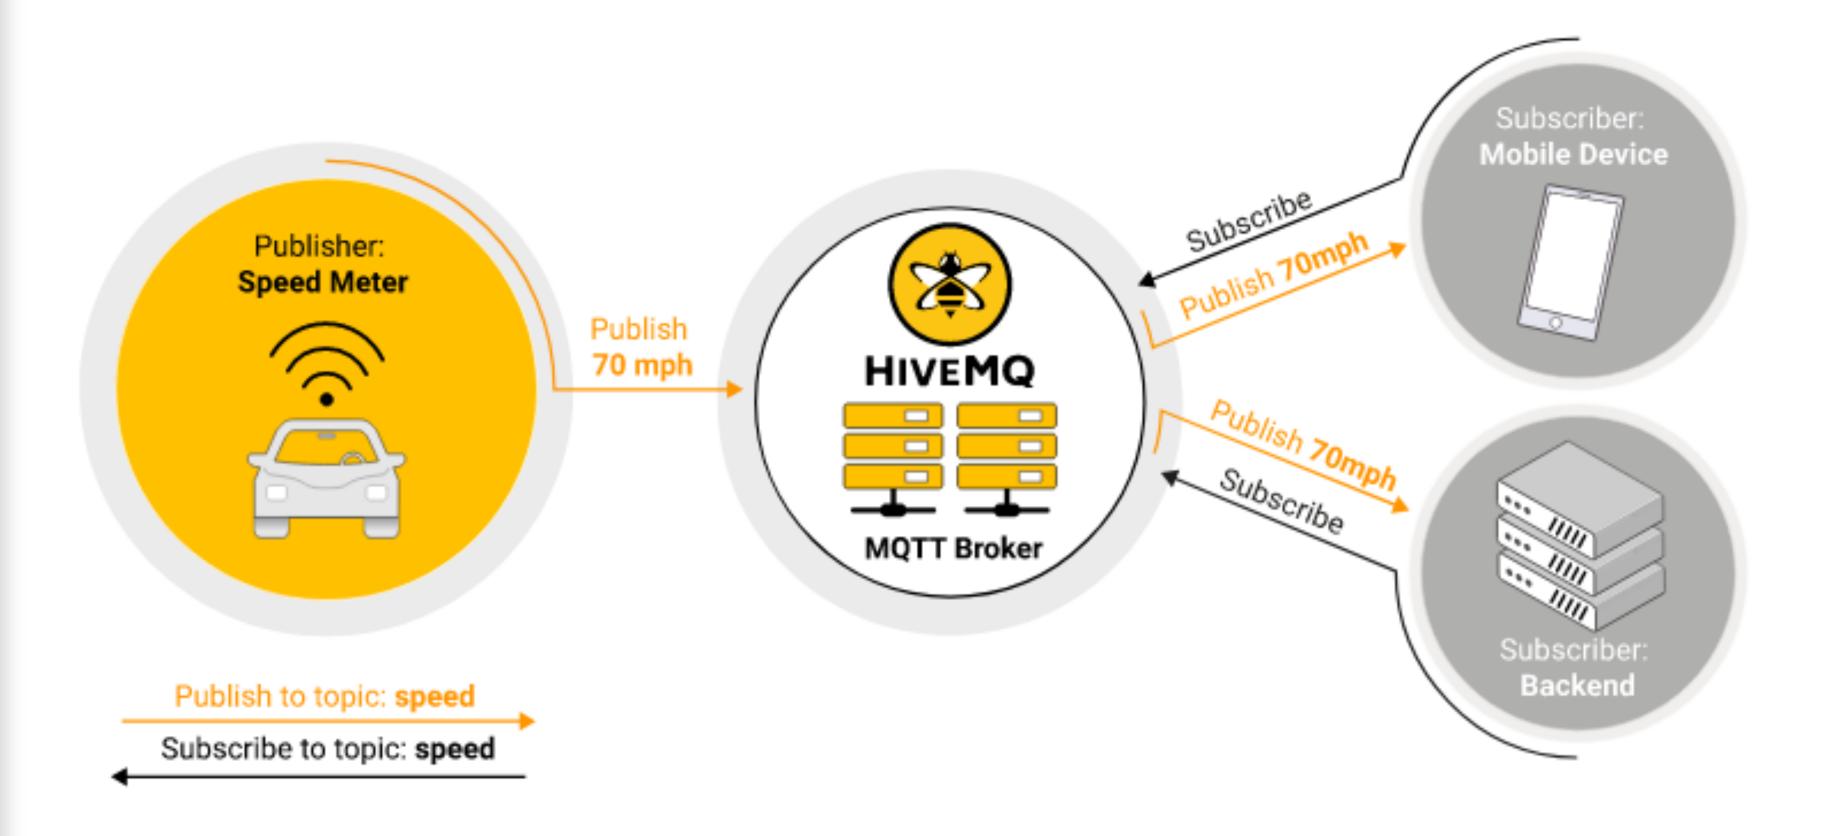


Image source: https://www.hivemq.com/blog/mqtt-essentials-part2-publish-subscribe/



MQTT: differences from Message Queues

In a Queue

- 1. Messages are stored until picked by a client (consumer)
 - → A message is always received by a client
- 2. A message is only consumed by one client
- 3. Queues are named and created explicitly
 - → messages can only be published & consumed after the queue is created

In a Topic (message "subject")

- 1. Messages are transient (i.e. not stored by the server, except for durable subscriptions)
 - → A messages may not be received by any client if no prior subscriptions are made
- 2. A message is received by all topic subscribers
- 3. Topics can be created on the fly
- → topics are created as the messages are published or consumers subscribe



MQTT Client

- Any device that runs an MQTT library and connects to an MQTT Broker over a network (TCP/IP stack)
 - E.g. A small resource-constraint device connected to a Broker via wireless network
 - E.g. A classic computer running an MQTT Client GUI for testing purposes e.g.
 MQTT Spy
- Publisher and/or subscriber
- Libraries for e.g., Android, Arduino, C, C++, C#, Go, iOS, Java, JavaScript, .NET.
 - Complete list: https://github.com/mqtt/mqtt.org/wiki/libraries
 - Eclipse Paho project open-source MQTT clients: https://www.eclipse.org/paho/



MQTT Broker

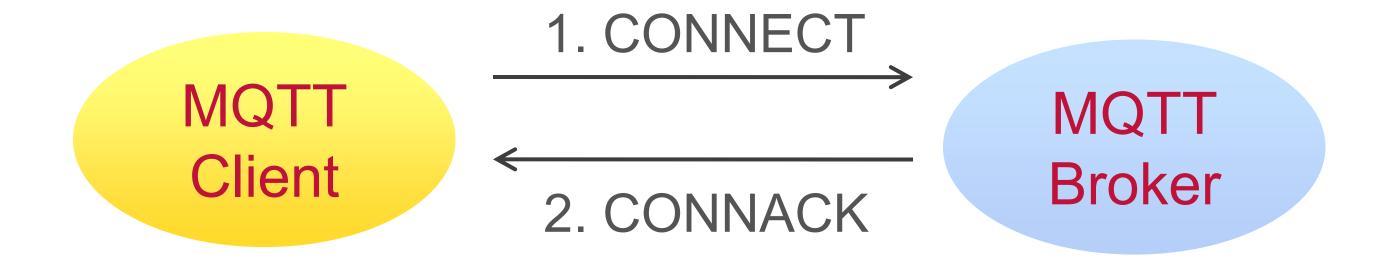
- Provides support for
 - Receiving messages
 - Filtering messages
 - Determining the subscribers of each message
 - Sending the message to these subscribers
 - Manage session data for clients with persistent sessions (e.g. subscriptions, missed messages)
 - Client authorisation and authentication
- Provider-specific properties
 - Scalability (e.g. millions of concurrent clients)
 - Failure-resistance
 - Integration into backend systems
 - Monitoring, ...

- Implementations
 - Mosquitto https://mosquitto.org/
 - Free , open-source (most popular)
 - Windows & Linux
 - Eclipse Paho MQTT
 - Free
 - HiveMQ http://www.hivemq.com/
 - Commercial
 - JoramMQ http://jorammq.com
 - Commercial
- Usage
 - Own locally installed Broker / Server
 - Cloud-based Server or Virtual Server
 - Shared Server Application



MQTT Connection

- MQTT: based on TCP/IP
 - Both Client and Broker need to have a TCP/IP stack
- Initiate Connection:
 - 1. Client sends Broker a CONNECT message (including KeepAlive parameter)
 - 2. Broker responds with CONNACK message & status code
 - → Broker keeps connection open until Client sends disconnect command or connection breaks (i.e. if Client fails to send PINGREQ before max delay, set via KeepAlive param.)





CONNECT message -> packet contents

- Client ID
 - Unique per client & broker (e.g. "Client-A") duplicate IDs cause disconnection of older Clients
 - Can be empty for stateless clients (must use true "clean session" flag)
- Clean session (flag) clean session = T : 新会话, 昨平不保留信息
 - Persistent session: CleanSession = false
 - → broker stores all client's sessions & non-acknowledged messages (if QoS > 0)
- Username/Password (optional)
 - Authentication (messages should be encrypted, e.g. TSL)
 - Alternative: SSL certificates
- Will message (optional)
 - Notifies other clients when a client disconnects ungracefully (e.g. failure)
- Keep Alive
 - Longest period (seconds) with no communication between Client & Broker (e.g. PING request/reply)



CONNACK message -> package contents

- Session present (flag)
 - Does the Borker already have a persistent session for the client?
 - Clients can know whether they are already subscribed to topics
 - Always: SessionPresent = false IF CleanSession == true
- Connect return code
 - Was the connection attempt successful or not?

Return Code	Meaning
0	Connection accepted
1	Connection refused, unacceptable protocol version
2	Connection refused, identifier rejected
3	Connection refused, server unavailable
4	Connection refused, bad user name or password
5	Connection refused, not authorized



PUBLISH

- Clients can publish messages after connecting to a Broker
- Client is unaware of how many subscribers there are for a topic
- Each message must contain a topic to determine the corresponding subscribers
- Packet Contains:
 - packetID (e.g. 1234) only relevant if QoS > 0; use 0 value otherwise
 - topicName: hierarchically structured String (e.g. "work/myOffice/temperature")
 - QoS: quality of service levels (i.e. 0, 1 or 2)
 - retainFlag: does the Broker save this message as the topic's last valid value?
 - payload: message content in byte format (client may use text, binary, xml, json, ...)
 - dupFlag: is this a duplicate message? (e.g. lost connection between publish & ack., QoS>0)



SUBSCRIBE

- To receive messages published on a topic, clients must subscribe to that topic
 - Clients only receive messages published to a topic after they subscribe to that topic
- SUBSCRIBE message (packet contents):
 - packetld (e.g. 1234)
 - List of subscriptions
 - qos<n>
 - topic<n>: subscription topic can contain wildchars
 - E.g. qos1,topic1, qos2,topic2,....
 - E.g. 1,"topic/1",0,"topic/2",....



SUBACK

- Broker acknowledges a client's subscription
- If Successful, the client receives all messages published on that topic
- Package content:
 - packageld: the same as for the Subscribe message
 - List of return codes: for each topic from Subscribe, in the same order

Return Code	Meaning
0	Success - Maximum QoS 0
1	Success - Maximum QoS 1
2	Success - Maximum QoS 2
128	Failure



UNSUBSCRIBE

- Deletes a Client's indicated subscriptions with the Broker
- Packet contents
 - packetld
 - List of topics: to unsubscribe from (e.g. "topic/1","/topic/2",...)

UNSUBACK

- Broker confirms unsubscription
- Pakcet contents
 - packetld: same as unsubscribe message



Topic

UTF-8 String, hierarchically structured by a forward slash "/" (topic level separator)

```
• E.g. "home/livingroom/light-sensor/1"

Topic Topic
Level Level

Topic
Level
Separator
```

- The Broker uses topics to filter messages to connected clients (subscribed)
- Clients do not need to create a topic before publishing or subscribing to it
- Wildcards can be used to subscribe to topics (not to publish)
 Wildcard (matched by any string)
 - Single-level: home/+/light-sensor/1
 - Multi-level: home/livingroom/#



System Topics

- Beginning with "\$"
 - Often used for the internal statistics of the MQTT Broker
 - Common topic: \$SYS/#
- Not included when subscribing to wildcard topic "#"
- E.g.,
 - \$SYS/broker/clients/connected
 \$SYS/broker/clients/disconnected
 \$SYS/broker/clients/total
 \$SYS/broker/messages/sent
 \$SYS/broker/uptime



Topics – Best Practices

- Do not use a leading forward slash: wasted topic level
- Avoid spaces in a topic: harder to read & debug
- Use concise & precise topic names: minimise network load
 → MQTTv5 uses numeric aliases for topic names
- Use only ASCII characters (printable): easier to debug
- Don't subscribe to all topics (#): scalability issues
- Consider future extensibility: how may your topics change?
- Use specific topics, rather than generic ones: e.g. 1 topic per sensor type



Quality of Service (QoS)

- Agreement between a message sender and a message receiver
 - Hence between Publisher (P) and Broker (B) and/or between Broker (B) and Subscriber (S)
 - End-to-end QoS (Publisher to Subscriber) involves both parts (P → B & B → S)
- Defines the guarantees of message delivery
- 3 QoS levels in MQTT:
 - 0: at most once

 - 2: exactly once



- Concerns two sides of message delivery
 - From Publisher to Broker: QoS defined by the publisher (in PUBLISH message)
 - From Broker to Subscriber: QoS defined by subscriber (in SUBSCRIBE message)
 - \rightarrow if Subscriber QoS is lower than Publisher QoS, the broker uses the lower QoS



QoS 0: deliver at most once

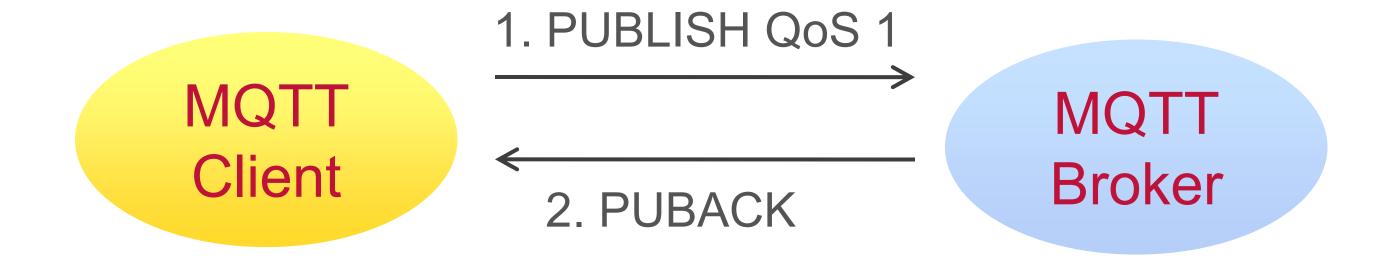
- Best effort delivery no guarantee, "fire & forget", same QoS as TCP/IP
- The recipient does not acknowledge message reception
- The message is *not* stored & resent by the sender





QoS 1: deliver at least once

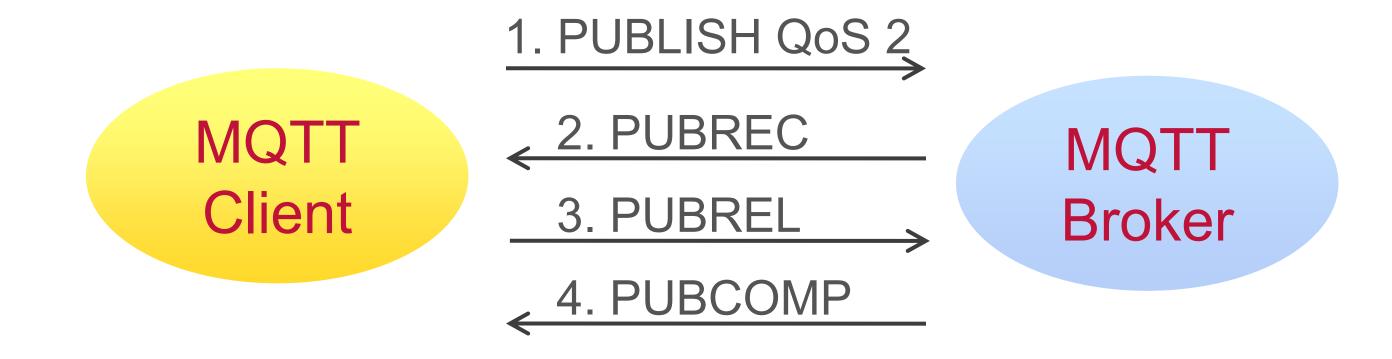
- Guarantees that the message is delivered at least once to the receiver
- The sender stores the message until it gets a PUBACK packet from the receiver
 - Uses the packetId to match the published packet to the acknowledgement
- The sender resends the packet if it does not receive a PUBACK within a delay;
 or if the connection was lost before the PUBACK arrived
 - Sets the duplicate flag DUP (not processed by broker or subscriber in QoS 1)
- A message may be delivered multiple times





QoS 2: deliver exactly once

- Guarantees that a message is received once and only once by intended recipients
 - Via 2x request/reply exchanges (4 part handshake) between sender & receiver
 - Uses the packetId of the initial message (e.g. PUBLISH)
- The sender resends the packet if it does not receive a PUBREC within a delay
 - Sets the duplicate flag DUP
- The receiver discards duplicates and sends a PUBCOMP message after receiving a PUBREL message
 - Sender discards all duplicates upon receiving PUBCOMP message
- Safest & slowest QoS level





QoS Best Practices

- How to select the "correct" QoS level? → Application & use-case dependent
- QoS 0
 - Mostly stable network connection (e.g. wired)
 - Acceptable to lose some messages (e.g. frequent measurements)
 - No need to queue messages (i.e. done for QoS 1 or 2; and persistent session)
- QoS 1
 - Need to receive every message & can handle duplicates
 - Cannot support the overhead of QoS 2 (i.e. faster message delivery with QoS 1)
- QoS 2
 - Critical to receive all messages exactly once (i.e. duplicate delivery can harm recipients)
 - Can support overheads and delays in message delivery



Persistent Session

- The Broker stores information for an offline client; available when client reconnects
 - Session existence (even if no subscription)
 - Client's subscriptions
 - All messages with QoS 1 or 2 that the client has not yet completely confirmed
 - All new messages with QoS 1 or 2 that the client missed when offline
- Clients can demand a persistent session when connecting to the Broker
 - cleanSession == true → non-persistent session;
 - → all info is lost if the client disconnects
 - cleanSession == false → persistent session
 - → all info is stored until the client requests a cleanSession
- Clients with persistent sessions must also store information (used upon reconnection)
 - All messages with QoS 1 or 2 that are not yet confirmed by the Broker



Best Practices – Persistent or Clean Session

Persistent session

- The Client must receive all messages from a topic, even if offline sometimes
- The Client has limited resources; the Broker can re-establish communication quickly
- The Client needs to resume all messages with QoS 1 & 2 after reconnect

Clean session

- The Client is only a publisher (not a subscriber)
- The Client does not need to get messages that it misses when offline



Retained Messages

- A message published with flag retain = true
- The Broker stores the last retained message and its QoS for a topic
 - A single retained message (i.e. the last one) is stored for each topic
- A Client that subscribes to a topic receives the retained message from that topic
 - Allows subscribed Clients to get a status update immediately (last known good value)
- Each new retained message replaces the old one on the same topic



To delete a retained message, send a retained message with zero payload to the targeted topic



Last Will & Testament (LWT)

- Unreliable networks, device batteries, ... > ungraceful disconnects of Clients
 - Clients can specify their LWT when connecting to the Broker
 - CONNECT message: lastWillTopic, lastWillQoS, lastWillMessage, lastWillRetain
- LWT allows to notify other Clients about a Client's ungraceful disconnect
 - MQTT message with a topic, retain flag, QoS and payload
- The Broker stores the LWT message until the Client disconnects ungracefully; then sends the LWT to all Clients subscribed to that topic
 - Broker detects IO error or network failure, no communication within Keep Alive period, connection closed without DISCONNECT message, protocol error
- If the Client disconnects gracefully (i.e. DISCONNECT message) then the Broker discards the LWT message
 - MQTTv5: Client may send DISCONNECT message & ask that the WILL be delivered



Keep Alive

- MQTT relies on Transmission Control Protocol (TCP)
 - Reliable, ordered, error-checked packet transfer over the Internet
- If one communication end fails (e.g. Mobile networks) -> half-open connection
 - The functioning end keeps sending messages and waits for acknowledgement
- Keep Alive: assess if a connection is still open (between a Client and a Broker)
- A Client specifies the Keep Alive interval upon connection to the Broker
 - Maximum interval with no communication (client-dependent: min. 0, max. 18h12min15 sec)
 - Client must send at least a PINGREQ packet if it has nothing else to communicate
 - The Broker responds with PINGRESP packet to confirm its availability to the Client
 - Broker disconnects Client if it receives no message within 1.5*KeepAlive interval



MQTT over Web Sockets

- Send & receive MQTT messages to a browser
 - Any browser that supports WebSockets can be an MQTT Client
 - Needs
 - JavaScript library that enables MQTT over WebSockets e.g. https://eclipse.org/paho/clients/js/
 - Broker that supports MQTT over WebSockets

WebSocket

- Network protocol: bi-directional, ordered, lossless comm. between a Browser & a Web Server
- Standardised in 2011; supported by all modern browsers
- Based on TCP

MQTT over WebSocket

- MQTT message is encapsulated within one or several WebSocket frames
- Browsers do not yet support Socket API (direct TCP connections)
- May use secure WebSockets via Transport Layer Security (TSL) to encrypt messages



JoramMQ



- Developed by: ScalAgent D.T., Grenoble, France
- Commercial: based on OW2/Joram 5.17, open-source broker (LGPL)
- Java implementation
- http://jorammq.com JoramMQ



QUESTIONS?

THANK YOU!



MQTT: Useful Information in a Nutshell

- Default MQTT port: 1883
- Used transport protocol: TCP/IP
 - Other protocols available for MQTT-SN (e.g. UDP, Zigbee)
- Data format for communication: binary
- Message encryption: not by default
- Each Client connected to a Broker must have a unique ClientId
- Is it possible to know the identity of a publisher Client? No
 - Unless the publisher includes its Id in the message payload
- Can I get a list of published topics? Not easily the Broker does not keep a list
 - Topics aren't permanent if a topic has no subscribers the Broker discards messages published on this topic



MQTT Resources

- MQTT essentials: https://www.hivemq.com/mqtt-essentials/
 - Text & Video formats
- MQTT Beginners Guide: http://www.steves-internet-guide.com/mqtt/
 - Related MQTT tutorials: http://www.steves-internet-guide.com/category/mqtt/
- MQTT Protocol homepage: https://mqtt.org/
- Eclipse IoT: http://iot.eclipse.org/getting-started
- MQTT Brokers/servers:
 - https://en.wikipedia.org/wiki/Comparison_of_MQTT_implementations
 - http://www.steves-internet-guide.com/mqtt-hosting-brokers-and-servers/