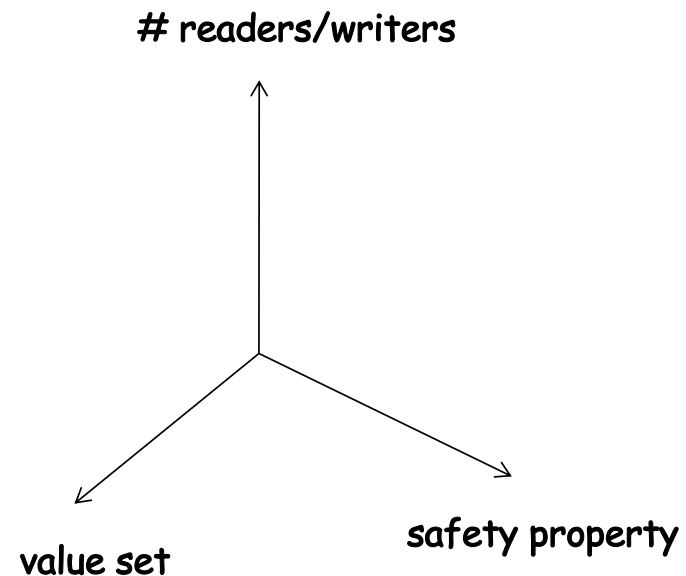


Atomic snapshot

SLR206, P2

The space of registers

- Nb of writers and readers: from 1W1R to NWNR
- Size of the value set: from binary to multi-valued
- Safety properties: safe, regular, atomic



All registers are (computationally) equivalent!

Transformations

From 1W1R binary safe to 1WNR multi-valued atomic

- I. From safe to regular (1W1R)
- II. From one-reader to multiple-reader (regular binary or multi-valued)
- III. From binary to multi-valued (1WNR regular)
- IV. From regular to atomic (1W1R)
- V. From 1W1R to 1WNR (multi-valued atomic)
- VI. From 1WNR to NWNR (multi-valued atomic)
- VII. From safe bit to atomic bit (optimal, coming later)

This class

- Atomic snapshot: reading multiple locations atomically
 - ✓ Write to one, read *all*

Atomic snapshot: sequential specification

- Each process p_i is provided with operations:

- ✓ $\text{update}_i(v)$, returns ok

- ✓ $\text{snapshot}_i()$, returns $[v_1, \dots, v_N]$

- In a **sequential** execution:

- For each $[v_1, \dots, v_N]$ returned by $\text{snapshot}_i()$,

- v_j ($j=1, \dots, N$) is the argument of the last $\text{update}_j(.)$

- (or the initial value if no such update)

Snapshot for free?

Code for process p_i :

initially:

shared 1WNR *atomic* register $R_i := 0$

upon snapshot()

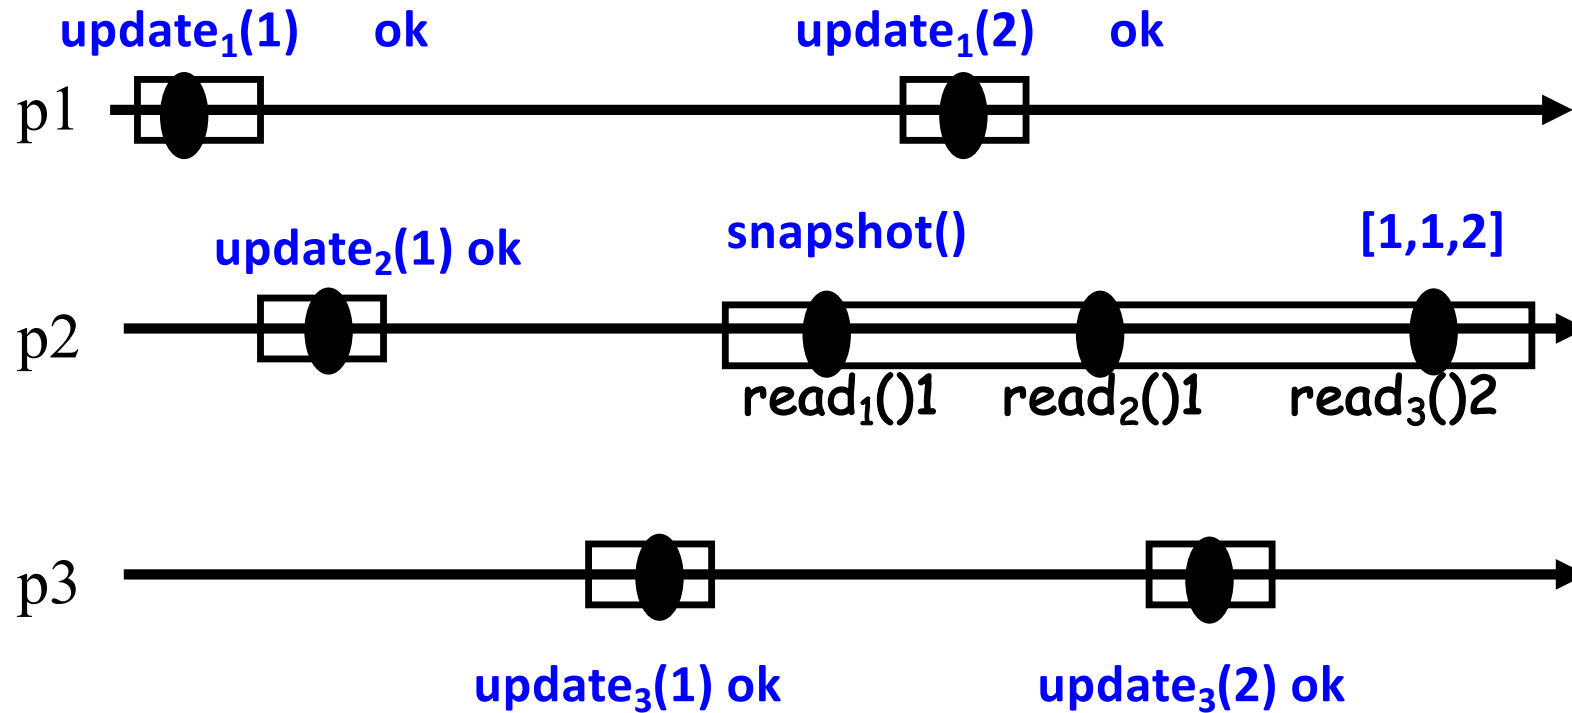
$[x_1, \dots, x_N] := \text{scan}(R_1, \dots, R_N)$ */*read R_1, \dots, R_N */*

return $[x_1, \dots, x_N]$

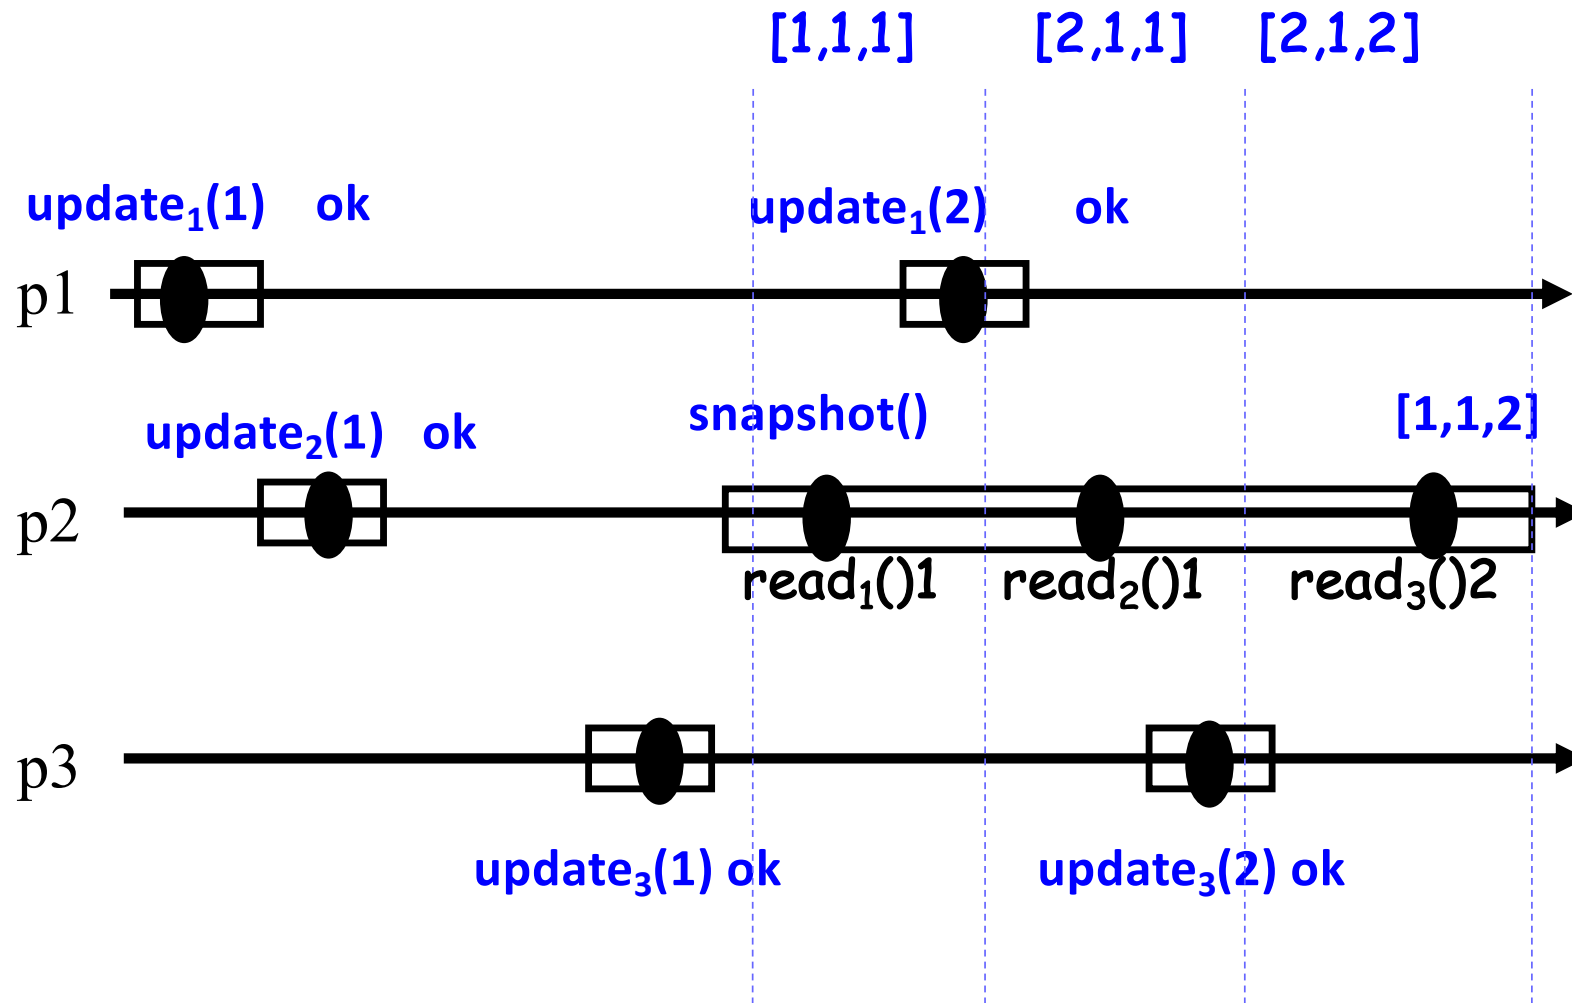
upon update_i(v)

$R_i.\text{write}(v)$

Snapshot for free?



Snapshot for free?



- What about 2 processes?
- What about **lock-free** snapshots?
 - ✓ At least one correct process **makes progress** (completes infinitely many operations)

Lock-free snapshot

Code for process p_i (all written values, including the initial one, are **unique**, e.g., equipped with a sequence number)

Initially:

shared 1W1R atomic register $R_i := 0$

upon snapshot()

$[x_1, \dots, x_N] := \text{scan}(R_1, \dots, R_N)$

repeat

$[y_1, \dots, y_N] := [x_1, \dots, x_N]$

$[x_1, \dots, x_N] := \text{scan}(R_1, \dots, R_N)$

until $[y_1, \dots, y_N] = [x_1, \dots, x_N]$

return $[x_1, \dots, x_N]$

upon update_i(v)


$R_i.\text{write}(v)$

Linearization

Assign a **linearization point** to each operation

- $\text{update}_i(v)$
 - ✓ $R_i.\text{write}(v)$ if present
 - ✓ Otherwise remove the op
- $\text{snapshot}_i()$
 - ✓ if complete – any point between identical scans
 - ✓ Otherwise remove the op

Build a **sequential history S** in the order of linearization points




Correctness: linearizability

S is legal: every $\text{snapshot}_i()$ returns the last written value for every p_j

Suppose not: $\text{snapshot}_i()$ returns $[x_1, \dots, x_N]$ and some x_j is not the the argument of the last $\text{update}_j(v)$ in S preceding $\text{snapshot}_i()$

Let C_1 and C_2 be two scans that returned $[x_1, \dots, x_N]$



Correctness: lock-freedom

An $\text{update}_i()$ operation is wait-free (returns in a finite number of steps)

Suppose process p_i executing $\text{snapshot}_i()$ eventually runs in isolation (no process takes steps concurrently)

- All scans received by p_i are distinct
- At least one process performs an update between
- There are only finitely many processes \Rightarrow at least one process executes infinitely many updates

What if base registers are regular?

General case: helping?

What if an update interferes with a snapshot?

- Make the update do the work!

WF

upon snapshot()

$[x_1, \dots, x_N] := \text{scan}(R_1, \dots, R_N)$

$[y_1, \dots, y_N] := \text{scan}(R_1, \dots, R_N)$

if $[y_1, \dots, y_N] = [x_1, \dots, x_N]$ then

return $[x_1, \dots, x_N]$

else

let j be such that

$x_j \neq y_j$ and $y_j = (u, U)$

return U

upon update_i(v)

$S := \text{snapshot}()$

$R_i.\text{write}(v, S)$


目的是为了直接使用上次 update 的 snapshot.

If two scans
differ - some
update succeeded!

Would this work?

问题是不确定上一次 update 的 snapshot
在当前 snapshot 范围内

Not that easy!



General case: wait-free atomic snapshot

upon snapshot()

$[x_1, \dots, x_N] := \text{scan}(R_1, \dots, R_N)$

while true do

$[y_1, \dots, y_N] := [x_1, \dots, x_N]$

$[x_1, \dots, x_N] := \text{scan}(R_1, \dots, R_N)$

if $[y_1, \dots, y_N] = [x_1, \dots, x_N]$ then

return $[x_1, \dots, x_N]$

else if moved_j and $x_j \neq y_j$ then

let $x_j = (u, U)$

return U

for each j: $\text{moved}_j := \text{moved}_j \vee x_j \neq y_j$

upon update_i(v)

$S := \text{snapshot}()$

$R_i.\text{write}(v, S)$

If a process moved
twice: its last
snapshot is valid!

通过两个判断确定在当前 snapshot

范围内同一 process 发生了两

次 snapshot, 那么最新的一

次 snapshot 必发生在当前范围内
验证了 snapshot 的有效性

lock
free

wait free

Correctness: wait-freedom

Claim 1 Every operation (update or snapshot) returns in $O(N^2)$ steps (bounded wait-freedom)

snapshot: does not return after a scan if a concurrent process moved and no process moved twice

- At most $N-1$ concurrent processes, thus (pigeonhole), after N scans:
 - ✓ Either at least two consecutive identical scans
 - ✓ Or some process moved twice!


update: snapshot() + one more step

Correctness: linearization points

update_i(v): linearize at the $R_i.write(v, S)$

complete **snapshot()**

- If two identical scans: between the scans
- Otherwise, if returned U of p_j : at the linearization point of p_j 's snapshot



The linearization is:

- Legal: every snapshot operation returns the most recent value for each process
- Consistent with the real-time order: each linearization point is within the operation's interval
- Equivalent to the run (locally indistinguishable)

(Full proof in the lecture notes, Chapter 6)

Atomic snapshot: sequential specification

- Each process p_i is provided with operations:

- ✓ $\text{update}_i(v)$, returns ok

- ✓ $\text{snapshot}_i()$, returns $[v_1, \dots, v_N]$

- In a **sequential** execution:

- For each $[v_1, \dots, v_N]$ returned by $\text{snapshot}_i()$,

- v_j ($j=1, \dots, N$) is the argument of the last $\text{update}_j(.)$

- (or the initial value if no such update)

One-shot atomic snapshot (AS)

Each process p_i :
 $\text{update}_i(v_i)$
 $S_i := \text{snapshot}()$

$S_i = S_i[1], \dots, S_i[N]$
(one position per process)

Vectors S_i satisfy:

- **Self-inclusion**: for all i : v_i is in S_i
- **Containment**: for all i and j :
 S_i is subset of S_j or S_j is subset of S_i

$$v_i \in S_j \Leftrightarrow S_j[i] = v_i \quad \text{bottom} \downarrow$$
$$S_i \subseteq S_j \Leftrightarrow \forall k, S_i[k] \neq \perp \Rightarrow S_i[k] = S_j[k]$$

Quiz 5: atomic snapshots

1. Prove that **one-shot** atomic snapshot satisfies self-inclusion and containment:
 - ✓ **Self-inclusion**: for all i : v_i is in S_i
 - ✓ **Containment**: for all i and j : S_i is subset of S_j or S_j is subset of S_i
2. Show that the atomic snapshot is subject to the **ABA problem** (affecting correctness) in case the written values are **not unique**