

SLR206

PROJECT

Optimistic Lock-Based List-Based Set Implementations

KANG Jiale

XIE Shifeng

October 18, 2023

Contents

1	Hand-Over-Hand Algorithm	3
1.1	Implementation	3
1.2	Proof of Safety	6
1.3	Proof of Liveness	7
2	Performance Analysis	8
2.1	Fixed update ratio 10% and varying list size	8
2.2	Fixed list size 100 and varying update ratios	9
2.3	Fixed update ratio 10% and list size 1000	12
3	System Information	13

1 HAND-OVER-HAND ALGORITHM

1.1 Implementation

```
1 package linkedlists.lockbased;
2 import java.util.concurrent.locks.Lock;
3 import java.util.concurrent.locks.ReentrantLock;
4
5 import contention.abstractions.AbstractCompositionalIntSet;
6
7 public class HandOverHandListIntSet extends AbstractCompositionalIntSet
8 {
9     // sentinel nodes
10    private Node head;
11    private Node tail;
12
13    public HandOverHandListIntSet(){
14        head = new Node(Integer.MIN_VALUE);
15        tail = new Node(Integer.MAX_VALUE);
16        head.next = tail;
17    }
18
19    /*
20     * Insert
21     *
22     * @see contention.abstractions.CompositionalIntSet#addInt(int)
23     */
24    @Override
25    public boolean addInt(int item){
26        head.lock();
27        Node pred=head;
28        Node curr=pred.next;
29        try {
30            curr.lock();
31            try {
32                while (curr.key < item){
33                    pred.unlock();
34                    pred = curr;
35                    curr = pred.next;
36                    curr.lock();
37                }

```

```
38     if (curr.key==item){
39         return false;}
40     Node node = new Node(item);
41     node.next=curr;
42     pred.next=node;
43     return true;
44 } finally{
45     curr.unlock();
46 }
47 }
48 finally{
49     pred.unlock();
50 }
51 }
52
53 /*
54  * Remove
55  *
56  * @see contention.abstractions.CompositionalIntSet#removeInt(int)
57  */
58 @Override
59 public boolean removeInt(int item){
60     head.lock();
61     Node pred=head;
62     Node curr=pred.next;
63     try {
64         curr.lock();
65         try {
66             while (curr.key < item){
67                 pred.unlock();
68                 pred = curr;
69                 curr = pred.next;
70                 curr.lock();
71             }
72             if (curr.key==item){
73                 pred.next=curr.next;
74                 return true;}
75             return false;
76         } finally{
77             curr.unlock();
78         } }
79     finally{
```

```
80     pred.unlock();
81 }
82 }
83
84 /*
85  * Contains
86  *
87  * @see contention.abstractions.CompositionalIntSet#containsInt(int
88  *)
89  */
90 @Override
91 public boolean containsInt(int item){
92     head.lock();
93     Node pred=head;
94     Node curr=pred.next;
95     try {
96         curr.lock();
97         try {
98             while (curr.key < item){
99                 pred.unlock();
100                 pred = curr;
101                 curr = pred.next;
102                 curr.lock();
103             }
104             return (curr.key==item);
105         } finally{
106             curr.unlock();
107         }
108     }
109     finally{
110         pred.unlock();
111     }
112 }
113
114 private class Node {
115
116     Node(int item) {
117         key = item;
118         next = null;
119         lock = new ReentrantLock();
120     }
121
122     public void lock() {
```

```
121     this.lock.lock();
122 }
123
124 public void unlock() {
125     this.lock.unlock();
126 }
127
128 public int key;
129 public Node next;
130 private final Lock lock;
131 }
132
133 @Override
134 public void clear() {
135     head = new Node(Integer.MIN_VALUE);
136     head.next = new Node(Integer.MAX_VALUE);
137 }
138
139 /**
140  * Non atomic and thread-unsafe
141  */
142 @Override
143 public int size() {
144     int count = 0;
145
146     Node curr = head.next;
147     while (curr.key != Integer.MAX_VALUE) {
148         curr = curr.next;
149         count++;
150     }
151     return count;
152 }
153 }
```

1.2 Proof of Safety

In the case of Hand-Over-Hand Algorithm, the primary safety property is mutual exclusion: only one thread can access a critical section at a time.

This algorithm uses a fine-grained locking strategy where each node in the linked list is associated with a lock. When a thread wants to access a node in the list, it acquires the lock associated with that node. If another thread holds the lock, the first thread must wait until the lock is released. Because of this locking strategy, when one thread holds the lock for a

particular node, no other thread can simultaneously hold the lock for the same node.

1.3 Proof of Liveness

Liveness of the Hand-Over-Hand Algorithm means that threads eventually make progress and do not get stuck in a waiting state.

When a thread has completed its operations on a node, it releases the lock associated with that node. This ensures that locks are not held indefinitely. When a thread releases the lock and the node, it would go to the critical section, and another thread holds that node and lock it. This shows that every threads make progress.

2 PERFORMANCE ANALYSIS

2.1 Fixed update ratio 10% and varying list size

We fixed the update ratio to 10% and changed the list size in {100, 1000, 10000}, changed number of threads in {1, 4, 6, 8, 10, 12}.

As shown in figure 1, Coarse Grained looks more stronger under the condition of less List Size and single thread. We propose an hypothesis that when using one thread, Coarse Grained Algorithm no needs to lock, while facing multiple threads, it locks the majority.

And when list size is smaller, it release the lock faster because of less operations in each list.




Figure 1: Coarse Grained

In the figure 2, Hand-Over-Hand Algorithm seems less affected by the changing of list size. But for the condition of single thread and less list size, it shows the best. The reason we proposed for this result is similar to the Coarse Grained Algorithm, which are using less locks and waiting less for other threads.

On the contrary, Lazy Linked Algorithm works better under the condition of multiple threads, as shown in figure 3. This is because Lazy Linked Algorithm is a non-blocking synchronization strategy. Non-blocking algorithms ensure that every process is correct.

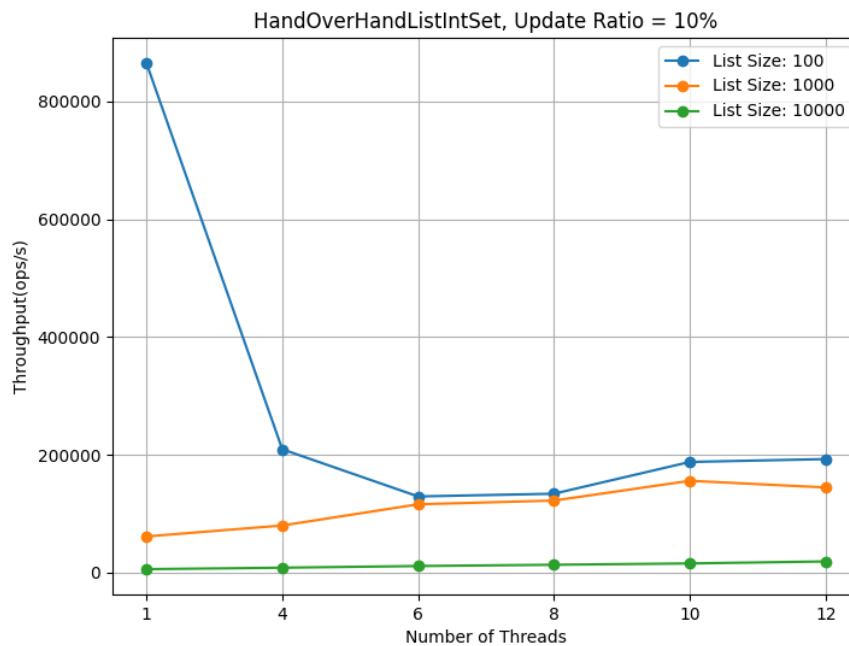


Figure 2: Hand-Over-Hand

Therefore, it is more efficient when using multiple threads.

2.2 Fixed list size 100 and varying update ratios

We fixed the list size to 100 and changed the update ratios in $\{0, 10\%, 100\%\}$, changed number of threads in $\{1, 4, 6, 8, 10, 12\}$.

For Coarse Grained Algorithm, as shown in figure 4, throughput varying less with update ratio. And single thread seems the most efficient condition. In the case of multiple threads, there is much room for improvement of the algorithm.

For Hand-Over-Hand Algorithm shown in figure 5, we proposed that throughput and update ratio is nearly independent (update ratio doesn't make big difference). The only reason decide the throughput is the number of threads.

For Lazy Linked Algorithm, update ratio affects a lot to the throughput, as shown in figure 6. Less update ratio will lead to more efficiency. And normally it works better when there are multiple threads.




Figure 3: Lazy Linked

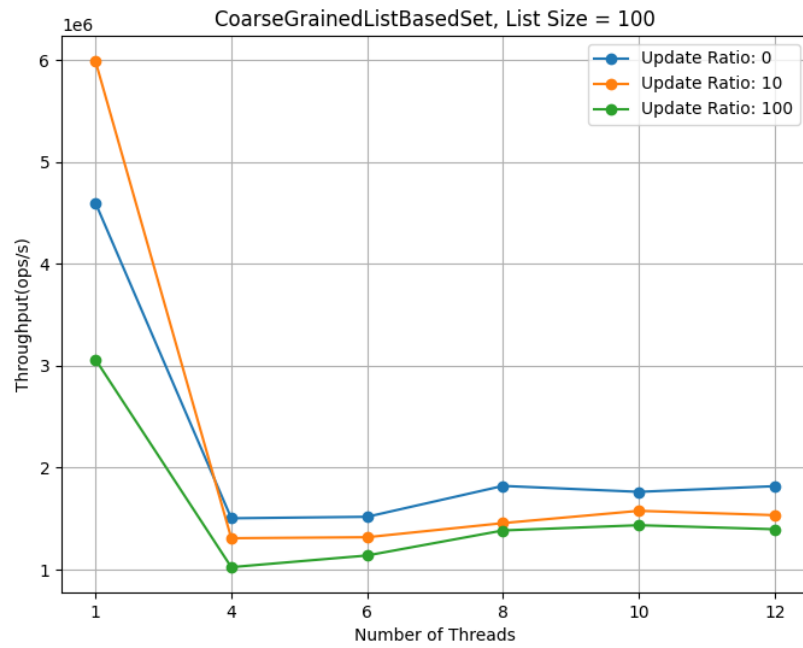


Figure 4: Coarse Grained

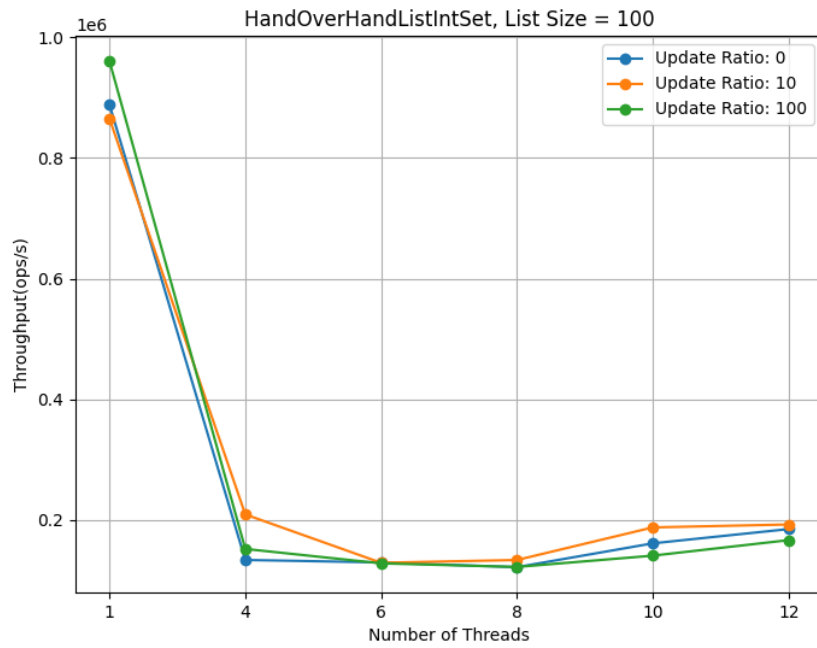


Figure 5: Hand-Over-Hand




Figure 6: Lazy Linked

2.3 Fixed update ratio 10% and list size 1000

We fixed the update ratio to 10% and list size to 1000, changed number of threads in {1, 4, 6, 8, 10, 12}.

Apparently, as figure 7 described, Lazy Linked Algorithm is the best algorithm among these three and it goes better when we increase threads, proving that it gets a big benefit with concurrency operations. Hand-Over-Hand Algorithm also rises throughput with the increase of the threads but do not change as large as Lazy Linked Algorithm. While Coarse Grained Algorithm goes decrease with the change of threads number.




Figure 7: Three Algorithms

3 SYSTEM INFORMATION

All the test are implemented in a remote server **@lamedell14.enst.fr**. Use command `lscpu` to check system CPU information.

```

1 Architecture:                x86_64
2   CPU op-mode(s):            32-bit, 64-bit
3   Address sizes:              43 bits physical, 48 bits virtual
4   Byte Order:                 Little Endian
5   CPU(s):                     128
6   On-line CPU(s) list:       0-127
7   Vendor ID:                  AuthenticAMD
8   Model name:                 AMD EPYC 7542 32-Core Processor
9   CPU family:                 23
10  Model:                       49
11  Thread(s) per core:         2
12  Core(s) per socket:         32
13  Socket(s):                   2
14  Stepping:                     0
15  BogomIPS:                    5789.50
16  Flags:                       fpu vme de pse tsc msr pae mce cx8 apic sep
                                mtrr pge mca cmov pat pse36 clflush mmx fxsr sse sse2 ht syscall nx
                                mmxext fxsr_opt pdpe1gb rdtscp lm constant_tsc rep_good nopl
                                nonstop_tsc cpuid extd_apicid aperfmperf rapl pni pclmulqdq monitor
                                ssse3 fma cx16 sse4_1 sse4_2 movbe popcnt aes xsave avx f16c rdrand
                                lahf_lm cmp_legacy svm extapic cr8_legacy abm sse4a misalignsse 3
                                dnowprefetch osvw ibs skinit wdt tce topoext perfctr_core perfctr_nb
                                bpext perfctr_llc mwaitx cpb cat_l3 cdp_l3 hw_pstate ssbd mba ibrs
                                ibpb stibp vmmcall fsgsbase bmi1 avx2 smep bmi2 cqm rdt_a rdseed adx
                                smap clflushopt clwb sha_ni xsaveopt xsavec xgetbv1 xsaves cqm_llc
                                cqm_occup_llc cqm_mbm_total cqm_mbm_local clzero irperf xsaveerptr
                                rdpru wbnoinvd amd_ppin arat npt lbrv svm_lock nrip_save tsc_scale
                                vmcb_clean flushbyasid decodeassists pausefilter pfthreshold avic
                                v_vmsave_vmload vgif v_spec_ctrl umip rdpid overflow_recov succor
                                smca sme sev sev_es
17 Virtualization features:
18   Virtualization:             AMD-V
19   Caches (sum of all):
20     L1d:                       2 MiB (64 instances)
21     L1i:                       2 MiB (64 instances)
22     L2:                        32 MiB (64 instances)
23     L3:                       256 MiB (16 instances)
24   NUMA:

```

```

25  NUMA node(s):                2
26  NUMA node0 CPU(s):          0-31,64-95
27  NUMA node1 CPU(s):          32-63,96-127
28  Vulnerabilities:
29  Itlb multihit:               Not affected
30  L1tf:                        Not affected
31  Mds:                          Not affected
32  Meltdown:                    Not affected
33  Mmio stale data:             Not affected
34  Retbleed:                     Mitigation; untrained return thunk; SMT
    enabled with ST
35                                IBP protection
36  Spec store bypass:           Mitigation; Speculative Store Bypass disabled
    via prctl
37                                and seccomp
38  Spectre v1:                   Mitigation; usercopy/swapgs barriers and
    __user pointer
39                                sanitization
40  Spectre v2:                   Mitigation; Retpolines, IBPB conditional,
    STIBP always-
41                                on, RSB filling, PBRSE-eIBRS Not affected
42  Srbds:                        Not affected
43  Tsx async abort:             Not affected

```