

Obstruction-free Consensus and Paxos

Jiale KANG & Yuanjie ZHAO

Overview

Our simulation of Synod algorithm:

- Export one operation $propose(v)$, $v \in \{0, 1\}$
- When a process invokes $propose(v)$, it will either *decide* or *abort*
- One process may invoke $propose(v)$ multiple times
- f faulty processes, with probability α to crash for each operation
- Leader-election mechanism: non-leader process stops proposing after t_{le}

Proof of correctness

- Validity

- Every decided value is a proposed value.
 - Obvious since the algorithm will not generate new values

- Agreement

- No two processes decide differently.
 - Suppose two processes i and j : p_i decides v with ballot b which is the lowest among all processes, p_j send an impose message with ballot b' and value v' where b' is the lowest ballot number higher than b attached to an impose message
 - There exists two majorities of processes, one for the impose phase of p_i and one for the read phase of p_j
 - Let p_k be a process in both majorities. p_k must have received p_i 's impose request before p_j 's read request in order not to abort. Thus, p_i must have received a gather response from p_k containing b and v . Thus by previous assumption of the ballot number of p_j , we can prove that $v=v'$, which means no two processes decide differently

Proof of correctness

- Obstruction-free termination

- If a correct process proposes, it eventually decides or aborts.
 - correct process receives either ABORT or GATHER
 - when receives GATHER from majority
 - sends IMPOSE and waits to receive ACKs
 - receives ACKs from majority, DECIDE
- If a correct process decides, no correct process aborts infinitely often.
 - before decides, send DECIDE to all
 - others when receive DECIDE, decide
- If there is a time after which exactly one correct process p proposes a value sufficiently many times, p eventually decides.
 - at some point, the proposal ballot b will be the greatest number in the system
 - cannot go to ABORT and eventually decides

Synod OFCons I

Code of every process p_i :

Initially:

```
ballot:=i-n; proposal:=nil; readballot:=0; imposeballot:=i-n;  
estimate:= nil; states:=[nil,0]n
```

upon propose(v)

```
proposal := v; ballot:=ballot + n; states:=[nil,0]n  
send [READ, ballot] to all
```

upon receive [READ,ballot'] from p_j

```
if readballot > ballot' or imposeballot > ballot' then  
  send [ABORT, ballot'] to  $p_j$   
else  
  readballot:=ballot'  
  send [GATHER, ballot', imposeballot, estimate] to  $p_j$ 
```

upon receive [ABORT, ballot] from some process

```
return abort
```

Synod OFCons II

```
upon receive [GATHER, ballot, estballot, est] from pj
    states[pj] := [est, estballot]

upon #states ≥ majority //collected a majority of responses
    if ∃ states[pk] = [est, estballot] with estballot > 0 then
        select states[pk] = [est, estballot] with highest estballot
        proposal := est //choose a potentially decided value
        states := [nil, 0]n
        send [IMPOSE, ballot, proposal] to all

upon receive [IMPOSE, ballot', v] from pj
    if readballot > ballot' or imposeballot > ballot' then
        send [ABORT, ballot'] to pj
    else
        estimate := v; imposeballot := ballot'
        send [ACK, ballot'] to pj
```

Synod OFCons III

upon received [ACK, ballot] from majority


send [DECIDE, proposal] to all

upon receive [DECIDE, v]


send [DECIDE, v] to all

return [decide, v]


Performance analysis




Performance analysis



Performance analysis



Performance analysis



Conclusion

- Larger systems use more time to complete consensus;
- Smaller systems show lower and stable consensus times;
 - less sensitive to the number of leader elections and crash probabilities
- Larger systems are more sensitive to t_{le} but still keeps stable for the variations of α .